


Distributed Systems  
*MIEEC, Fall 2018*

# Logical clocks

**Luis Almeida / Pedro Souto**  
**DEEC – University of Porto, Portugal**

# Temporal order

- Many applications do not require actual time
  - **Order** is enough
    - e.g., causality, versions control, distributed storage...
- Enforcing order in a distributed system is not trivial...
  - e.g., due to concurrent paths




# Precedence between events

- Certain events have **clear precedence** relationship
  - **Happened-before** relationship ( $\rightarrow$ ) proposed by L. Lamport
    - **HB1** (***events in the same process** always precede each other*):  
If  $e$  and  $e'$  are events occurring in one process in such order then  $e \rightarrow e'$
    - **HB2** (***transmission of messages** must always precede reception*):  
If  $e$  is the transmission of a message and  $e'$  its reception then  $e \rightarrow e'$
    - **HB3: (**transitivity**)**  
if  $e \rightarrow e'$  and  $e' \rightarrow e''$  then  $e \rightarrow e''$
  - The HB relationship is **partial** since there may exist pairs of events to which it does not apply  $\rightarrow$  **Concurrent events**

$$\neg (a \rightarrow e) \wedge \neg (e \rightarrow a) \Leftrightarrow a \parallel e$$

# Precedence between events

- The HB relationship captures the control flow




$a \rightarrow b ; b \rightarrow c ; c \rightarrow d ; d \rightarrow f ; e \rightarrow f$

$e \parallel (a, b, c, d)$

# Lamport clocks and timestamps

- Beyond the HB relationship **Lamport** also proposed:
  - A **logical clock** per process that counts its relevant events
    - *The Lamport clock*
  - The association of **timestamps** of such clock to **events**
    - *The Lamport timestamps -  $L(e)$*
- Rules for a **Lamport clock  $L_i$  of process  $p_i$** 
  - **LC1** (**increments** on each local event):
    - $L_i = L_i + 1$  (just before executing an event)
  - **LC2** (**synchronizes** on message reception):
    - Every message  $m$  carries the timestamp of the sender ( $m, t=L_i$ )
    - The receiver  $p_j$  adjusts its clock to enforce HB:  $L_j = \max(L_j, t) + 1$

# Lamport clocks and timestamps



- The relationship between **HB** and **Lamport timestamps**

$$e \rightarrow e' \Rightarrow L(e) < L(e')$$

$$L(e) < L(e') \quad ???$$

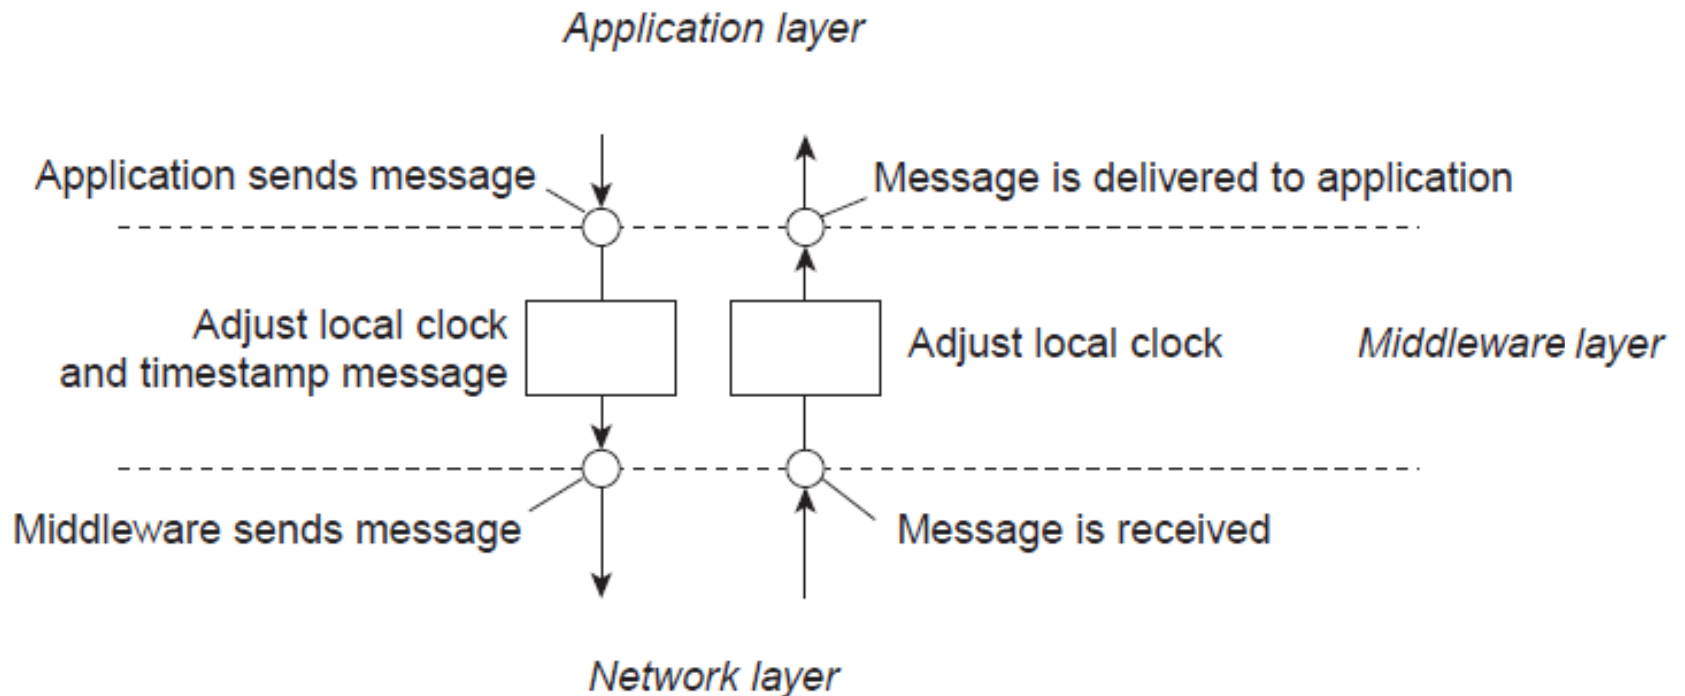
$$L(e) \geq L(e') \Rightarrow \neg (e \rightarrow e')$$

Cannot distinguish  
**precedence** from  
**concurrency**!

Distinguishes **not precedence** only!

# Implementing Lamport clocks

- **Middleware layer**
  - Separates message **reception** from **delivery**
  - Adjusts local clock




# Vector clocks

- Vector with *vision* of **all clocks** in all processes
  - Independently proposed by Mattern and Fidge in 1988
- Process  $p_i$  has vector  $V_i$ 
  - **VC1:** position  $i$  counts **local events**
    - $V_i[i] = V_i[i] + 1$  (just before executing an event)
  - **VC2:** **timestamps** are the whole **vector**
    - $t = V_i$  (and sent with every message)
  - **VC3:** position  $j$  updated when receiving a **message from  $p_j$** 
    - $V_i[j] = \max(V_i[j], t[j])$
- Comparing vectors
  - $V < V'$  iff  $\forall_j V[j] \leq V'[j] \wedge \exists_i V[i] < V'[i]$



# Vector clocks




## .The relationship between **HB** and **Vector clocks**

$e \rightarrow e' \Leftrightarrow V(e) < V(e')$  → Distinguishes **precedence** from **concurrency**!

$\neg (V(e) < V(e')) \wedge \neg (V(e') < V(e)) \Leftrightarrow e \parallel e'$  → **Concurrent** events

# Vector clocks



(from wikipedia)

# Further readings

- Tanenbaum and van Stenn, Distributed Systems
  - Section 6.2: Logical clocks
- The wikipedia has interesting high level info
  - For a quick reference
- Very interesting sequence of post in the Basho Blog
  - <http://basho.com/posts/technical/vector-clocks-revisited/>
- Very interesting presentation on Version Vectors
  - <https://www.youtube.com/watch?v=3SWSw3mKApM>